

# No-Show Appointment Data - Project 1\_v2

prepared by Nicole Wambui on 9th August 2022 (resubmitted version)

## Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
  - [Research Question 1](#)
  - [Research Question 2](#)
  - [Research Question 3](#)
  - [Research Question 4](#)
- [Conclusions](#)
- [Recommendations & Further Research](#)

## Introduction

### Motivation

The purpose of this project is to investigate the *No Show Appointment* dataset in order to get some useful insights about health in Brazil. I personally chose to focus on health because it has been a prevalent topic over the last few years due to the ongoing pandemic.

### Research Questions

There are the questions I will be exploring with the data:

1. Are certain age groups more likely to miss their appointments?
2. Do SMS prompts increase the likelihood of patients attending their appointments?
3. Are "*Bosla Familia*" (scholarship) beneficiaries more or less likely to attend appointments?
4. What are the highest risk factors among patients?

## Data Wrangling

### Loading the data

The data was already provided on this [link \(https://www.google.com/url?q=https://d17h27t6h515a5.cloudfront.net/topher/2017/October/59dd2e9a\\_noshowappointments-kagglev2-may-2016/noshowappointments-kagglev2-may-2016.csv&sa=D&ust=1532469042118000\)](https://www.google.com/url?q=https://d17h27t6h515a5.cloudfront.net/topher/2017/October/59dd2e9a_noshowappointments-kagglev2-may-2016/noshowappointments-kagglev2-may-2016.csv&sa=D&ust=1532469042118000) and downloaded in .csv format.

In [1]:

```
import pandas as pd

appts_df = pd.read_csv('NoShowApp.csv')
```

## Inspecting the data

In this step, I'd like to just look at the data, see what it looks like and note down all the errors that need to be taken care of before the data can be cleaned to suit my analysis.

In [2]:

```
appts_df.head(10)
```

Out[2]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood
0	2.987250e+13	5642903	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA
1	5.589978e+14	5642503	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA
2	4.262962e+12	5642549	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA
3	8.679512e+11	5642828	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI
4	8.841186e+12	5642494	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA
5	9.598513e+13	5626772	F	2016-04-27T08:36:51Z	2016-04-29T00:00:00Z	76	REPÚBLICA
6	7.336882e+14	5630279	F	2016-04-27T15:05:12Z	2016-04-29T00:00:00Z	23	GOIABEIRAS
7	3.449833e+12	5630575	F	2016-04-27T15:39:58Z	2016-04-29T00:00:00Z	39	GOIABEIRAS
8	5.639473e+13	5638447	F	2016-04-29T08:02:16Z	2016-04-29T00:00:00Z	21	ANDORINHAS
9	7.812456e+13	5629123	F	2016-04-27T12:48:25Z	2016-04-29T00:00:00Z	19	CONQUISTA

In [3]:

```
#check for null values in the data. We use the all() function to avoid reprinting the entire
appts_df.isnull().all()
```

Out[3]:

```
PatientId      False
AppointmentID  False
Gender         False
ScheduledDay   False
AppointmentDay  False
Age            False
Neighbourhood  False
Scholarship    False
Hipertension   False
Diabetes       False
Alcoholism     False
Handcap        False
SMS_received   False
No-show        False
dtype: bool
```

In [4]:

```
#We cannot assume that the number of patients and appointments is equal, some patients may
print(appts_df['PatientId'].nunique())
print(appts_df['AppointmentID'].nunique())
```

```
62299
110527
```

In [5]:

```
appts_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110527 entries, 0 to 110526
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   PatientId             110527 non-null float64
 1   AppointmentID          110527 non-null int64
 2   Gender                 110527 non-null object
 3   ScheduledDay           110527 non-null object
 4   AppointmentDay         110527 non-null object
 5   Age                   110527 non-null int64
 6   Neighbourhood          110527 non-null object
 7   Scholarship            110527 non-null int64
 8   Hipertension           110527 non-null int64
 9   Diabetes               110527 non-null int64
10   Alcoholism             110527 non-null int64
11   Handcap                110527 non-null int64
12   SMS_received           110527 non-null int64
13   No-show                110527 non-null object
dtypes: float64(1), int64(8), object(5)
memory usage: 11.8+ MB
```

In [6]:

```
#Because my quesitons are focused on individuals, I will have to drop duplicate values of t
appts_df['PatientId'].drop_duplicates(inplace=False)
appts_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110527 entries, 0 to 110526
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   PatientId             110527 non-null  float64
1   AppointmentID         110527 non-null  int64
2   Gender                110527 non-null  object
3   ScheduledDay          110527 non-null  object
4   AppointmentDay        110527 non-null  object
5   Age                   110527 non-null  int64
6   Neighbourhood         110527 non-null  object
7   Scholarship           110527 non-null  int64
8   Hipertension          110527 non-null  int64
9   Diabetes              110527 non-null  int64
10  Alcoholism            110527 non-null  int64
11  Handcap               110527 non-null  int64
12  SMS_received          110527 non-null  int64
13  No-show               110527 non-null  object
dtypes: float64(1), int64(8), object(5)
memory usage: 11.8+ MB
```

This is the first roadblock. If there are only 62299 unique patient IDs, yet after dropping duplicates the **PatientId** column remains the same, and the entire observaion numbers remain the same. It makes more sense to work with the data as it is, and only drop the patientid column to avoid confusion.

The **ScheduledDay** and **AppointmentDay** are listed as objects, we can also go ahead and see if we can parse them as dates, and any important insights we can get.

In [7]:

```
appts_df['ScheduledDay'] = pd.to_datetime(appts_df['ScheduledDay'], format='%Y-%m-%dT%H:%M:
appts_df['AppointmentDay'] = pd.to_datetime(appts_df['AppointmentDay'], format='%Y-%m-%dT%H
appts_df.head(2)
```

Out[7]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood
0	2.987250e+13	5642903	F	2016-04-29 18:38:08	2016-04-29	62	JARDIM DA PENHA
1	5.589978e+14	5642503	M	2016-04-29 16:08:27	2016-04-29	56	JARDIM DA PENHA

Here, we can see that the appointment day has no actual time value. Calculating the difference between the scheduling and the actual appointment may have been an interesting insight, but it is not really useful as it is. Maybe if we could figure out the day of the week, we may get some other insights, but that is out of scope for this project.

We can look at the Neighbourhoods and see if we can determine if certain neighbourhoods get more patients.

In [8]:

```
appts_df.Neighbourhood.value_counts()
```

Out[8]:

JARDIM CAMBURI	7717
MARIA ORTIZ	5805
RESISTÊNCIA	4431
JARDIM DA PENHA	3877
ITARARÉ	3514
...	
ILHA DO BOI	35
ILHA DO FRADE	10
AEROPORTO	8
ILHAS OCEÂNICAS DE TRINDADE	2
PARQUE INDUSTRIAL	1

Name: Neighbourhood, Length: 81, dtype: int64

\*81 Neighbourhoods would be too many to map out and so we can also safely exclude this from our dataset.

Based on the questions I have, there are a few variables I will not need such as : *patientid*, *appointmentid*, *scheduled* and *appointmentday*, and *neighbourhood*. We can focus on the variables I actually need:

In [9]:

```
appts_df.Age.describe()
```

Out[9]:

count	110527.000000
mean	37.088874
std	23.110205
min	-1.000000
25%	18.000000
50%	37.000000
75%	55.000000
max	115.000000

Name: Age, dtype: float64

In [10]:

```
#We cannot have an age that is less than 0, so let's look for values that are 0 and less than
appts_df.query('Age<=0')
```

Out[10]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbour
59	7.184428e+13	5638545	F	2016-04-29 08:08:43	2016-04-29	0	CONQU
63	2.366233e+14	5628286	M	2016-04-27 10:46:12	2016-04-29	0	SÃO BENE
64	1.885174e+14	5616082	M	2016-04-25 13:28:21	2016-04-29	0	ILHA CAIE
65	2.718818e+14	5628321	M	2016-04-27 10:48:50	2016-04-29	0	CONQU
67	8.647128e+13	5639264	F	2016-04-29 08:53:02	2016-04-29	0	PALE
...	...	...	...	...	...	...	...
110345	1.473952e+14	5702537	F	2016-05-16 12:30:58	2016-06-01	0	RESISTÊ
110346	5.577525e+12	5777724	M	2016-06-06 14:22:34	2016-06-08	0	RESISTÊ
110454	6.142460e+11	5772400	F	2016-06-03 15:18:44	2016-06-03	0	RESISTÊ
110460	4.321846e+13	5769545	F	2016-06-03 08:56:51	2016-06-03	0	RESISTÊ
110507	4.769462e+14	5786918	F	2016-06-08 09:04:18	2016-06-08	0	MARIA C

3540 rows × 14 columns



From here, we can see that there are 3540 different observations with children listed as 0 years, a good guess would be that they are infants. The best thing we can do here is include them, but remove the entire  $age = -1$  observation.

In [11]:

```
print(appts_df.shape) #check the number of rows
appts_df.drop(appts_df[appts_df['Age'] < 0].index, inplace = True)
print(appts_df.shape) #check the number of rows to ensure we deleted it
```

(110527, 14)

(110526, 14)

In [12]:

```
appts_df.Age.describe() #confirm the min value is now 0 and the data is valid
```

Out[12]:

```
count    110526.000000
mean      37.089219
std       23.110026
min        0.000000
25%       18.000000
50%       37.000000
75%       55.000000
max       115.000000
Name: Age, dtype: float64
```

Everything else is pretty much straight forward so we can go ahead and remove the other columns.

In [13]:

```
appts_df.drop(['PatientId', 'AppointmentID', 'ScheduledDay', 'AppointmentDay', 'Neighbourhood'])
appts_df.head(2)
```

Out[13]:

	Gender	Age	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap	SMS_received	N shc
0	F	62	0	1	0	0	0	0	1
1	M	56	0	0	0	0	0	0	1

Now we can rename the columns and save this as a new dataframe for exploratory analysis.

In [14]:

```
appts_df.rename(columns=lambda x: x.strip().lower().replace("-", "_"), inplace=True) #this
#there are also spelling errors with hypertension and handicap, we can fix that with a sim
appts_df.rename(columns={'hipertension': 'hypertension', 'handcap': 'handicap'}, inplace=True)
appts_df.head(1)
```

Out[14]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received	no_s
0	F	62	0	1	0	0	0	0	

In [15]:

```
#save as a data frame and export it.
appts_df.to_csv('NoShowApp_v2.csv', index=False)
```

# Exploratory Data Analysis

In this section, we will go ahead and answer the 5 research questions that we laid out in the beginning to get better insights on the data:

1. Do SMS prompts increase the likelihood of patients attending their appointments?
2. Are certain age groups more likely to miss their appointments?
3. Are "Bosla Familia" (scholarship) beneficiaries more or less likely to attend appointments?
4. What are the highest risk factors among patients?

In [16]:

```
#load the cleaned dataset as well as all the packages needed

import pandas as pd #no need to load it again, but I will do for good measure
import numpy as np
import seaborn as sns #this package helps us create a background for the visualisations
import matplotlib.pyplot as plt
%matplotlib inline
#including a space after the magic(%) sign results in an error

appts_df = pd.read_csv('NoShowApp_v2.csv')
appts_df.head(1)
```

Out[16]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received	no_s
0	F	62	0	1	0	0	0	0	

## Question 1: Do SMS prompts increase the likelihood of patients attending their appointments?

In [17]:

```
#we can create a dataframe using the function that will filter the results for us, this wil
sms_df = appts_df[(appts_df['sms_received'] == 1)]
```

In [18]:

```
sms_df.sms_received.value_counts() #confirm it is only showing results where the sms was re
```

Out[18]:

```
1    35482
Name: sms_received, dtype: int64
```




In [19]:

```
#we can now query the same data frame to get value counts for "no-shows", this will give us
print(sms_df['no_show'].value_counts()) #print this to know how many were no shows and how
no_shows_df1 = sms_df['no_show'].value_counts()
print((no_shows_df1[0]),(no_shows_df1[1])) #confirm we can print the output if we create a
```

```
No      25698
Yes      9784
Name: no_show, dtype: int64
25698 9784
```

In [21]:

```
colors = ['blue'] #not necessary at the moment, but it will be later
plt.subplots(figsize=(8, 5))
plt.bar(("Missed Appointment", "Attended Appointment"),[no_shows_df1[1], no_shows_df1[0]])
plt.title('Number of Appointments Missed where SMS was received')
plt.xlabel('Attendance')
plt.ylabel('Number of appointments');
```



9,784 appointments were missed and 25,698 were attended for those who received sms's.

We can now plot the number of people who didn't receive sms's and look at their attendance.

In [22]:

```
sms_df2 = appts_df[(appts_df['sms_received'] == 0)]
sms_df2.sms_received.value_counts()
```

Out[22]:

```
0      75044
Name: sms_received, dtype: int64
```

In [23]:

```
print(sms_df2['no_show'].value_counts())
no_shows_df2 = sms_df2['no_show'].value_counts()
print((no_shows_df2[0]),(no_shows_df2[1]))
```

```
No      62509
Yes     12535
Name: no_show, dtype: int64
62509 12535
```


In [24]:

```
#Let's create a handy function to make plotting easier
```

```
def plt_labels(a,b,c):
    plt.title(a)
    plt.xlabel(b)
    plt.ylabel(c)
```

In [25]:


```
colors = ['blue'] #not necessary at the moment, but it will be later
plt.subplots(figsize=(8, 5))
plt.bar(("Missed Appointment", "Attended Appointment"),[no_shows_df2[1], no_shows_df2[0]])
plt_labels('Number of Appointments Missed where SMS was not received','Attendance','Number
#the function lets us call it(plt_labels) instead of having to write out 3 similar lines fo
```



12, 535 appointments were missed whereas 62, 509 were attended despite not getting an sms.

In [26]:

```
colors = ['blue', 'orange'] #blue is for sms received and orange is for sms not received. We
plt.subplots(figsize=(8, 5))
plt.bar(("Missed + SMS", "Attended + SMS"),[no_shows_df1[1], no_shows_df1[0]])
plt.bar(("Missed", "Attended"),[no_shows_df2[1], no_shows_df2[0]])
plt_labels('Attendance of Appointment vs SMS', 'Attendance and SMS', 'Number of appointments')
```



#### Answer:

Receiving an sms does not improve the likelihood of attending an appointment. We can, however, suggest that sms prompts reduce the number of missed appointments.

## Question 2: Are certain age groups more likely to miss their appointments?

In order to answer this question, we need to create age categories before we can plot this.

In [27]:

```
#to create the age groups, we can either create set categories i.e 0-10,10-20 etc. or use the
appts_df.age.describe()
```

Out[27]:

```
count      110526.000000
mean         37.089219
std          23.110026
min           0.000000
25%          18.000000
50%          37.000000
75%          55.000000
max         115.000000
Name: age, dtype: float64
```

In [28]:

```
#we will use the min, 25thpercentile, median, 75thpercentile and max as our category boundaries
ages = (appts_df.age.describe()) #create a new dataframe
age_groups = [ages[3], ages[4], ages[5], ages[6], ages[7]] #create age group categories as
age_group_names = ['0-18 yrs', '19-37 yrs', '38-55 yrs', '56-115 yrs'] #create category names
```

In [29]:

```
appts_df['age_categories'] = pd.cut(appts_df['age'], age_groups, labels=age_group_names) #create
appts_df.head(7) #confirm the column was successfully created
```

Out[29]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received	no_show
0	F	62	0	1	0	0	0	0	
1	M	56	0	0	0	0	0	0	
2	F	62	0	0	0	0	0	0	
3	F	8	0	0	0	0	0	0	
4	F	56	0	1	1	0	0	0	
5	F	76	0	1	0	0	0	0	
6	F	23	0	0	0	0	0	0	

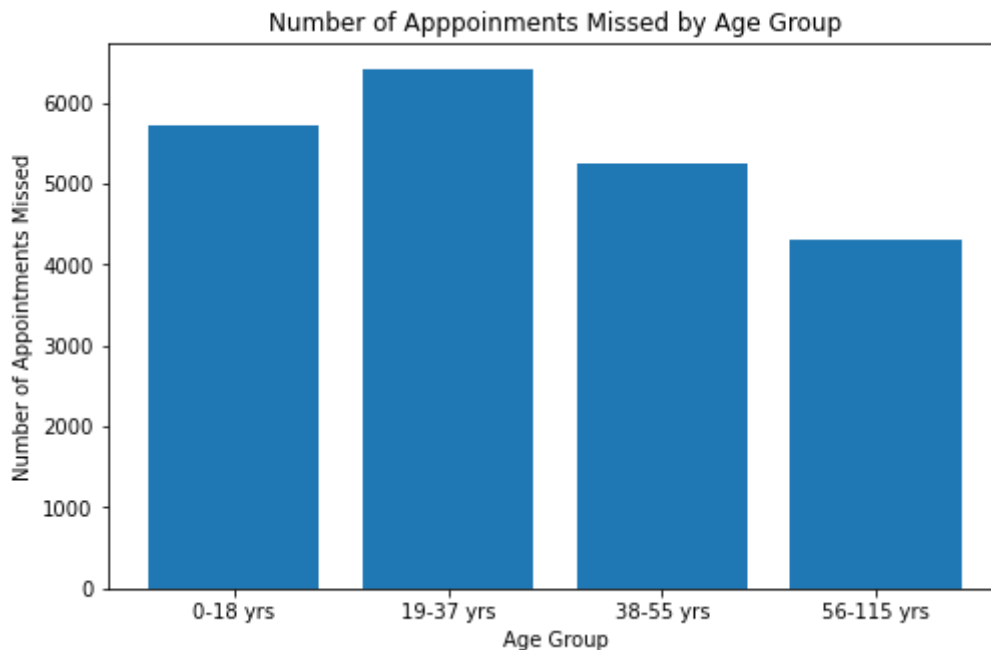
In [30]:

```
age_data = appts_df[(appts_df['no_show'] == "Yes")] #create a dataframe with filtered values
print(age_data['age_categories'].value_counts()) #print the numbers of the ages so that we
age_data = age_data['age_categories'].value_counts()
```

```
19-37 yrs      6414
0-18 yrs       5708
38-55 yrs      5255
56-115 yrs     4303
Name: age_categories, dtype: int64
```

In [31]:

```
#we are first of all interested in the general number of appointments for each
plt.subplots(figsize=(8, 5))
plt.bar(("0-18 yrs", "19-37 yrs", "38-55 yrs", "56-115 yrs"),[age_data[1], age_data[0], age_data[1], age_data[0]])
plt_labels('Number of Appointments Missed by Age Group', 'Age Group', 'Number of Appointments Missed')
```

**Answer:**

People within the age group of 19 to 37 years are more likely to miss an appointment, followed by people within the 0 to 18 years age group.

### Question 3: Are "*Bosla Familia*" (scholarship) beneficiaries more or less likely to attend appointments?

In [32]:

```
#we need to first filter the data set to look at people who are scholarship beneficiaries and then we can plot this all at once using the same method we used in the first research question.
schol_df1 = appts_df[(appts_df['scholarship'] == 1)]
schol_df2 = appts_df[(appts_df['scholarship'] == 0)] #create 2 different data frames focusing on scholarship beneficiaries and non-scholarship beneficiaries
```

In [33]:

```
print(schol_df1['no_show'].value_counts()) #print this to know how many were no shows and how many were not
appts_schol_1 = schol_df1['no_show'].value_counts()
print((appts_schol_1[0]),(appts_schol_1[1])) #print this to ensure the vectors work
```

```
No      8283
Yes     2578
Name: no_show, dtype: int64
8283 2578
```

In [34]:


```
#we'll do the same thing again for the non-beneficiaries (schol_df2)
print(schol_df2['no_show'].value_counts())
appts_schol_2 = schol_df2['no_show'].value_counts()
print((appts_schol_2[0]),(appts_schol_2[1]))
```

```
No      79924
Yes      19741
Name: no_show, dtype: int64
79924 19741
```

In [35]:

```
#we can now go ahead and plot these two side-by-side for a more efficient visualisation

colors = ['blue','orange'] #blue is for scholarship beneficiaries and orange is for sms non
plt.subplots(figsize=(8, 5))
plt.bar(("Missed+scholarship", "Attended+scholarship"),[appts_schol_1[1], appts_schol_1[0]])
plt.bar(("Missed", "Attended"),[appts_schol_2[1], appts_schol_2[0]])
plt_labels('Attendance of Appointment vs Scholarship Status','Attendance and Scholarship St
```



In [36]:

```
#there seem to be significantly more non-beneficiaries than beneficiaries. We can try to fi
schol_df = schol_df2.shape[0] - schol_df1.shape[0]
print(schol_df)
```

88804

Scholarship beneficiaries miss fewer appointments. However, there are far fewer beneficiaries of the scholarship than non-beneficiaries. There are exactly 88,804 more non-beneficiaries than beneficiaries.

In [37]:

```
#We can try to find the difference in missed vs. attended appointments for beneficiaries and  
beneficiaries_df = appts_schol_2[0] - appts_schol_2[1] #will calculate difference for attended  
nonbeneficiaries_df = appts_schol_1[0] - appts_schol_1[1] #will calculate difference for missed
```

In [38]:

```
print(beneficiaries_df)  
print(nonbeneficiaries_df)
```

60183

5705

### Answer:

When comparing those who missed their appointments between beneficiaries and non-beneficiaries the differences were:

- The difference between attended and missed appointments was 60,183 for non-beneficiaries.
- The difference between attended and missed appointments was 5,705 for beneficiaries.

We can suggest that a fewer number of beneficiaries miss appointments in comparison to non-beneficiaries.

## Question 4: What are the highest risk factors among patients?

This question serves as my univariate analysis and I will simply be exploring the 4 risk factors present in the dataset:

1. Hypertension
2. Diabetes
3. Alcoholism
4. Handicap

In [39]:

```
#we will first look at the handicap column because to have a better understanding of the data  
print(appts_df.hypertension.value_counts())  
print(appts_df.diabetes.value_counts())  
print(appts_df.alcoholism.value_counts())  
print(appts_df.handicap.value_counts())
```

```
0    88725  
1    21801  
Name: hypertension, dtype: int64  
0    102583  
1     7943  
Name: diabetes, dtype: int64  
0    107166  
1     3360  
Name: alcoholism, dtype: int64  
0    108285  
1     2042  
2       183  
3        13  
4         3  
Name: handicap, dtype: int64
```

We can see that the handicap column has 3 values, most likely related to the level/type of handicap. We need to handle this data slightly differently before we can plot it.

In [40]:

```
#we'll create data frames from the above code to simplify the plotting  
hypertension = appts_df.hypertension.value_counts() #Map the diseases to different variables  
diabetes = appts_df.diabetes.value_counts()  
alcoholism = appts_df.alcoholism.value_counts()
```

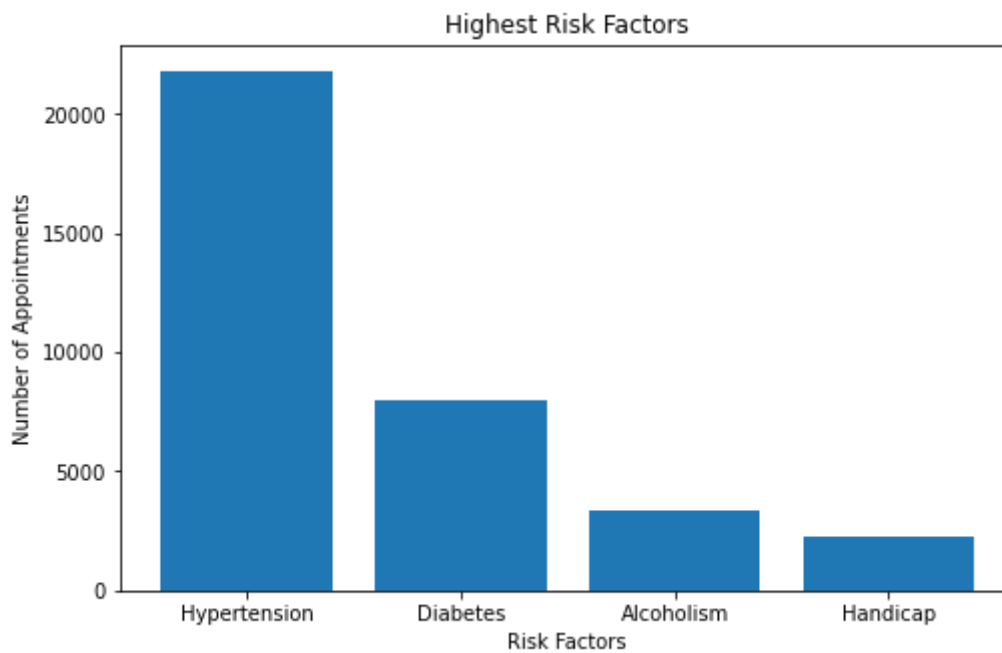
In [41]:

```
#we'll use the sum of the vectors as the full data frame for the handicap column i.e.  
ha = appts_df.handicap.value_counts()  
handicap = ha[1]+ha[2]+ha[3] #this data frame includes the added values of the last 3 unique values
```



In [42]:

```
color=['array'] #choose different colours for each column
plt.subplots(figsize=(8, 5))
plt.bar(("Hypertension", "Diabetes", "Alcoholism", "Handicap"), [hypertension[1], diabetes[1],
plt_labels('Highest Risk Factors', 'Risk Factors', 'Number of Appointments')
```



Here, we can see that **Hypertension** is the highest risk factor associated with appointments.

In [43]:

```
#we will now deep dive into hypertension and gain some statistics.

hypertension_df = appts_df[(appts_df['hypertension'] == 1)]
```

In [44]:

```
hypertension_df.head(5)
```

Out[44]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received	no_
0	F	62	0	1	0	0	0	0	
4	F	56	0	1	1	0	0	0	
5	F	76	0	1	0	0	0	0	
25	M	46	0	1	0	0	0	1	
26	F	45	0	1	0	0	0	0	

We can see something quite interesting here. One patient appears to have diabetes as well. We can investigate even further.

In [45]:

```
hypertension_df[(hypertension_df['diabetes'] == 1)]
```

Out[45]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received	
4	F	56	0	1	1	0	0	0	
36	F	63	0	1	1	0	0	0	
37	F	64	1	1	1	0	0	0	
44	F	78	0	1	1	0	0	0	
47	F	39	0	1	1	0	0	0	
...	...	...	...	...	...	...	...	...	...
110358	F	70	0	1	1	0	0	0	
110447	F	40	0	1	1	0	0	1	
110455	F	79	0	1	1	0	0	0	
110467	F	76	0	1	1	0	0	1	
110498	F	66	0	1	1	0	0	0	

6486 rows × 10 columns

In [46]:

```
hypertension_df[(hypertension_df['alcoholism'] == 1)]
```

Out[46]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received
46	M	58	0	1	0	1	0	1
186	M	66	0	1	0	1	0	0
381	F	54	0	1	0	1	0	0
384	M	63	0	1	0	1	0	0
746	F	56	0	1	0	1	0	0
...	...	...	...	...	...	...	...	...
109885	M	71	0	1	0	1	0	0
109911	M	56	0	1	1	1	0	0
110070	M	54	0	1	0	1	0	1
110167	M	50	0	1	1	1	0	0
110173	M	59	0	1	0	1	0	1

1327 rows × 10 columns



In [47]:

```
hypertension_df[(hypertension_df['handicap'] == 1)]
```

Out[47]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received
147	F	65	0	1	0	0	1	0
189	F	77	0	1	0	0	1	0
199	M	62	0	1	1	0	1	0
211	F	56	0	1	0	0	1	0
238	F	65	0	1	1	0	1	1
...	...	...	...	...	...	...	...	...
109743	M	28	0	1	0	0	1	1
109776	M	80	0	1	1	0	1	1
109790	M	66	0	1	0	0	1	0
109853	F	58	0	1	1	0	1	1
110024	M	63	1	1	0	0	1	0

883 rows × 10 columns



In [48]:

```
hypertension_df[(hypertension_df['alcoholism'] == 1)|(hypertension_df['diabetes'] == 1)|(hy
```

Out[48]:

	gender	age	scholarship	hypertension	diabetes	alcoholism	handicap	sms_received
4	F	56	0	1	1	0	0	0
36	F	63	0	1	1	0	0	0
37	F	64	1	1	1	0	0	0
44	F	78	0	1	1	0	0	0
46	M	58	0	1	0	1	0	1
...	...	...	...	...	...	...	...	...
110358	F	70	0	1	1	0	0	0
110447	F	40	0	1	1	0	0	1
110455	F	79	0	1	1	0	0	0
110467	F	76	0	1	1	0	0	1
110498	F	66	0	1	1	0	0	0

8098 rows × 10 columns

In [49]:

```
hypertension_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21801 entries, 0 to 110514
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   gender                 21801 non-null object  
1   age                    21801 non-null int64   
2   scholarship            21801 non-null int64   
3   hypertension           21801 non-null int64   
4   diabetes               21801 non-null int64   
5   alcoholism             21801 non-null int64   
6   handicap               21801 non-null int64   
7   sms_received           21801 non-null int64   
8   no_show                21801 non-null object  
9   age_categories         21801 non-null category
dtypes: category(1), int64(7), object(2)
memory usage: 1.7+ MB
```

In [50]:

```
(6846 + 1327 + 883) == 8098 #number of rows from this function hypertension_df[(hypertension_
```

Out[50]:

False

In [51]:

```
(6846 + 1327 + 883) #totaling the figures gotten from these codes hypertension_df[(hypertensi
```

Out[51]:

9056

In [59]:

```
#we'll now try to plot these comorbidities in a pie chart (6846-diabtes, 1327-alcoholism, 8
```

```
figures = np.array([21801, 6846, 1327, 883])
```


```
color=['array'] #choose different colours for each column
```

```
plt.subplots(figsize=(8, 8))
```

```
plt.pie(figures, labels = ["Hypertension-only", "Diabetes", "Alcoholism", "Handicap"])
```

```
plt.title('Hypertension Comorbidities'); #we won't call the plt_label function because pie
```

Hypertension Comorbidities



**Answer:** Hypertension is the most common risk factor.

Of the 21,801 people who have hypertension, 6486 have diabetes, 1327 suffer from alcoholism and 883 have a handicap. While the above outputs do not match, we can suggest that people who suffer from hypertension are likely to present other risk factors. This could explain why they have many appointments.

## Conclusions

From this dataset, it can be suggested that:

1. SMS prompts do not necessarily increase the likelihood of attending appointments, but there was a noticeable reduction in the number of appointments missed for those who had SMS prompts.
2. People between the ages of 19 to 37 are more likely to miss their appointments.
3. People who are beneficiaries of *Bosla Famila* are less likely to miss their appointments.
4. Hypertension is the most prevalent risk factor,
5. People with Hypertension and another risk factor are most likely to present with Diabetes.

## Recommendations for Further Analysis

It would be best to first state the limitations of this analysis. First of all, differences in gender and region were not accounted for. Secondly, there was no consideration for the date of the scheduling and the actual appointment. This was mostly because the preliminary analysis showed that the data had negative values when trying to calculate the difference in days. This was not explored any further due to time constraints.

It would be a worthy endeavour to investigate the following:

- Specific days or times where most appointments are missed.
- The number of *Bosla Famila* subscriptions per region.
- The correlation between age and risk factors.
- The relationship between risk factors and likelihood of missed appointments.